Machine Learning - 4105931

Lecture 3
Types of Learning

Chen-Kuo Chiang (江 振 國)
ckchiang@cs.ccu.edu.tw
中正大學 資訊工程學系

Types of Learning

- Learning with Different Output 1. Space $\mathcal{Y}$
- Learning with Different 2. Data Label y_n
- Learning with Different 3. Protocol $f \Rightarrow (\mathbf{x}_n, y_n)$
- Learning with Different 4. Input Space $\mathcal{X}$

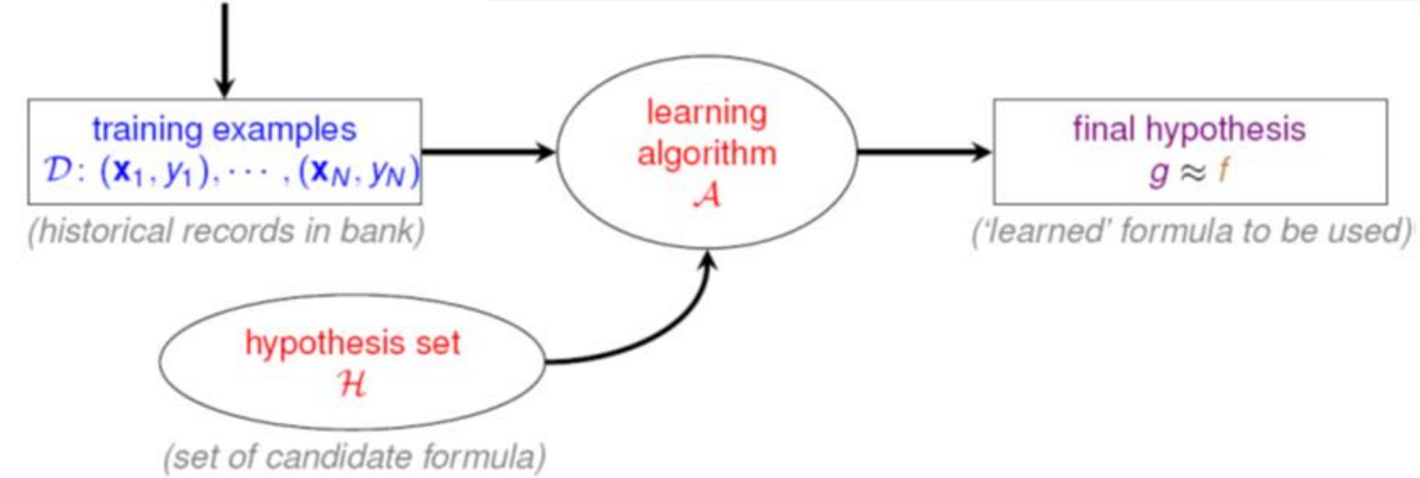

Learning with Different Output Space Y:

Credit Approval Problem Revisited

age	23 years
gender	female
annual salary	NTD 1,000,000
year in residence	1 year
year in job	0.5 year
current debt	200,000

credit? $\{$**no**(-1), **yes**$(+1)\}$

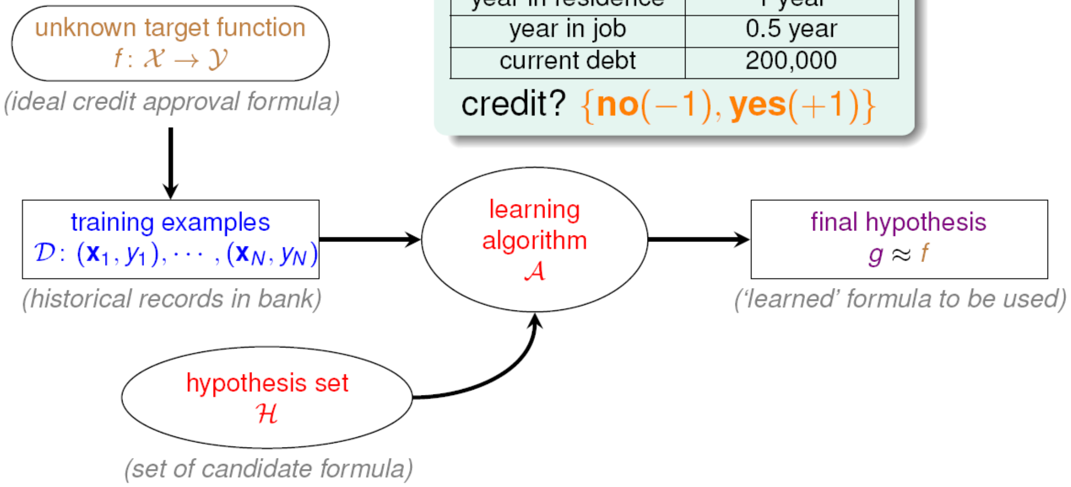

$\mathcal{Y} = \{-1, +1\}$: **binary classification**

More Binary Classification Problems

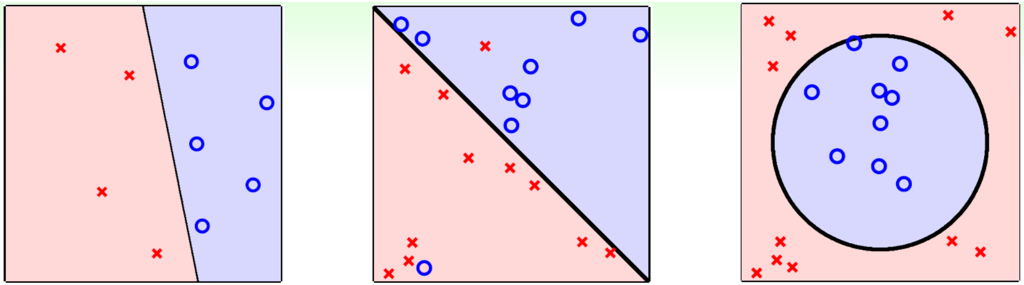

- credit approve/disapprove
- email spam/non-spam
- patient sick/not sick
- ad profitable/not profitable
- answer correct/incorrect (KDDCup 2010)

core and important problem with
many tools as **building block of other tools** 二元分類

Multiclass Classification: Coin Recognition Problem

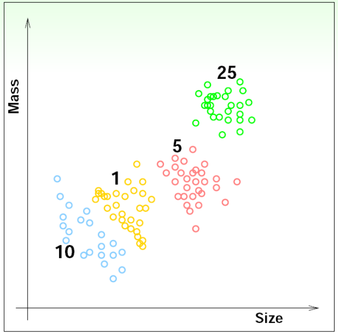

- classify US coins (1c, 5c, 10c, 25c) by (size, mass)
- $\mathcal{Y} = \{1c, 5c, 10c, 25c\}$, or $\mathcal{Y} = \{1, 2, \cdots, K\}$ **(abstractly)**
- binary classification: special case with $K = 2$

多類別

Other Multiclass Classification Problems

- written digits $\Rightarrow 0, 1, \cdots, 9$
- pictures $\Rightarrow$ apple, orange, strawberry
- emails $\Rightarrow$ spam, primary, social, promotion, update (Google)

many applications in practice,
especially for 'recognition'

Regression: Patient Recovery Prediction Problem

- binary classification: patient features $\Rightarrow$ sick or not 二元分類
- multiclass classification: patient features $\Rightarrow$ which type of cancer 多元分類
- regression: patient features $\Rightarrow$ **how many days before recovery** 預測
- $\mathcal{Y} = \mathbb{R}$ or $\mathcal{Y} = [\text{lower}, \text{upper}] \subset \mathbb{R}$ (bounded regression)
  —**deeply studied in statistics** 實數 區間預測 病人多久可以康復

Other Regression Problems

- company data $\Rightarrow$ stock price
- climate data $\Rightarrow$ temperature

also core and important with many 'statistical' tools as **building block of other tools**

Structured Learning: Sequence Tagging Problem

$$\underbrace{\mathit{I}}_{\mathit{pronoun}}\ \underbrace{\mathit{love}}_{\mathit{verb}}\ \underbrace{\mathit{ML}}_{\mathit{noun}}$$

- multiclass classification: word $\Rightarrow$ word class
- structured learning:
  sentence $\Rightarrow$ structure (class of each word)
- $\mathcal{Y} = \{PVN, PVP, NVN, PV, \cdots\}$, not including $VVVVV$
- huge multiclass classification problem (structure $\equiv$ hyperclass) **without ‘explicit’ class definition** 字母又可以構成詞

Other Structured Learning Problems

- protein data $\Rightarrow$ protein folding
- speech data $\Rightarrow$ speech parse tree

a fancy but complicated learning problem

Mini Summary

Learning with Different Output Space $\mathcal{Y}$

- **binary classification**: $\mathcal{Y} = \{-1, +1\}$
- multiclass classification: $\mathcal{Y} = \{1, 2, \cdots, K\}$
- **regression**: $\mathcal{Y} = \mathbb{R}$
- structured learning: $\mathcal{Y} =$ structures
- ... and a lot more!!

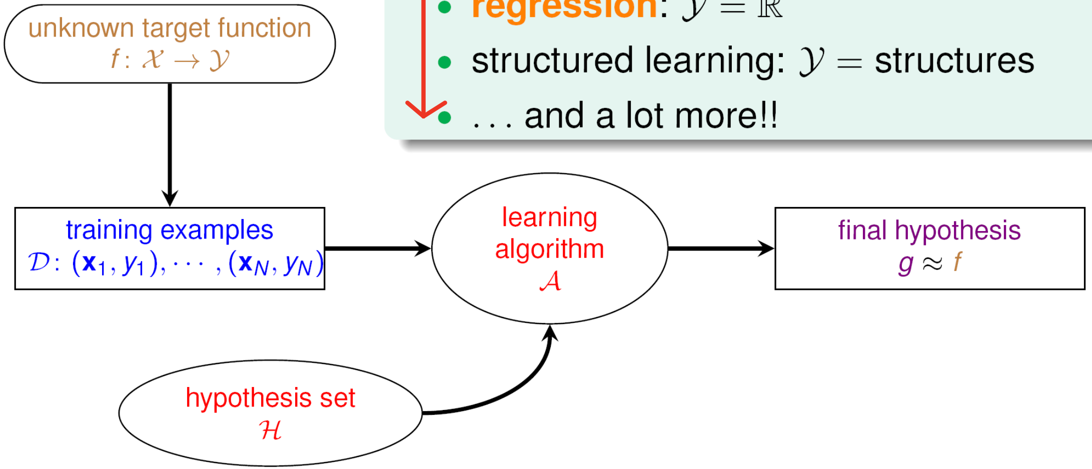

core tools: binary classification and regression

Fun Time

What is this learning problem?

The entrance system of the school gym, which does automatic face recognition based on machine learning, is built to charge four different groups of users differently: Staff, Student, Professor, Other. What type of learning problem best fits the need of the system?

1. binary classification
2. multiclass classification
3. regression
4. structured learning

Reference Answer: 2

There is an 'explicit' $\mathcal{Y}$ that contains four classes.

Learning with Different Data Label y_n

有答案

- Supervised: Coin Recognition Revisited

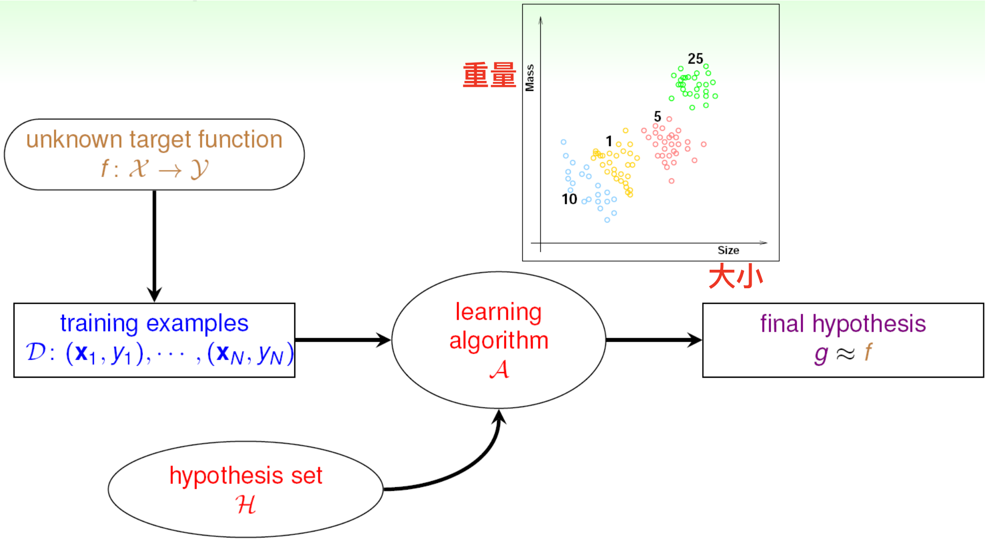

supervised learning:
every $\mathbf{x}_n$ **comes with corresponding** y_n

沒有答案

Unsupervised: Coin Recognition without y_n

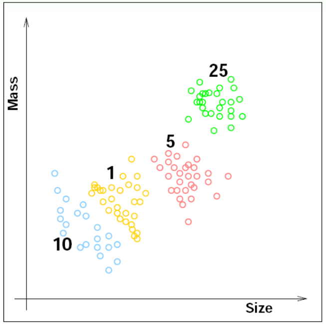

supervised multiclass classification

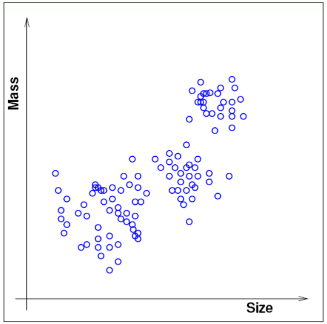

unsupervised multiclass classification

$\Longleftrightarrow$ '**clustering**'

分群

Other Clustering Problems

- articles $\Rightarrow$ topics
- consumer profiles $\Rightarrow$ consumer groups

clustering: a challenging but useful problem

Unsupervised: Learning without y_n

Other Unsupervised Learning Problems

- clustering: $\{\mathbf{x}_n\} \Rightarrow \text{cluster}(\mathbf{x})$
  ($\approx$ 'unsupervised multiclass classification')
  —i.e. articles $\Rightarrow$ topics
- **density estimation**: $\{\mathbf{x}_n\} \Rightarrow \text{density}(\mathbf{x})$
  ($\approx$ 'unsupervised bounded regression') 交通流量問題
  —i.e. traffic reports with location $\Rightarrow$ dangerous areas
- **outlier detection**: $\{\mathbf{x}_n\} \Rightarrow \text{unusual}(\mathbf{x})$ 外來者偵測
  ($\approx$ extreme 'unsupervised binary classification')
  —i.e. Internet logs $\Rightarrow$ intrusion alert
- . . . and a lot more!!

unsupervised learning: diverse, with possibly very different performance goals

Semi-supervised: Coin Recognition with Some y_n

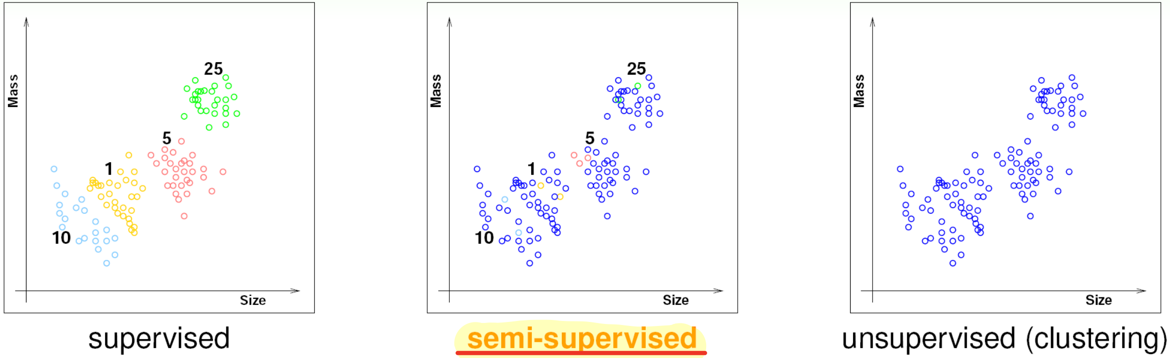

Other Semi-supervised Learning Problems

- face images with a few labeled $\Rightarrow$ face identifier (Facebook)
- medicine data with a few labeled $\Rightarrow$ medicine effect predictor

semi-supervised learning: leverage unlabeled data to avoid 'expensive' labeling

Reinforcement Learning

Teach Your Dog: Say 'Sit Down'

The dog pees on the ground.

BAD DOG. THAT'S A VERY WRONG ACTION.

- cannot easily show the dog that y_n = sit when $\mathbf{x}_n =$ 'sit down'
- but can 'punish' to say $\tilde{y}_n$ = pee is wrong

Teach Your Dog: Say 'Sit Down'

The dog sits down. 坐下會給獎勵

Good Dog. Let me give you some cookies.

- still cannot show y_n = sit when $\mathbf{x}_n =$ 'sit down'
- but can 'reward' to say $\tilde{y}_n$ = sit is good

Reinforcement Learning

- A ‘very different’ but natural way of learning

Other Reinforcement Learning Problems Using ($\mathbf{x}, \tilde{y}$, goodness)

- (customer, ad choice, ad click earning) $\Rightarrow$ $\boxed{\text{ad system}}$
- (cards, strategy, winning amount) $\Rightarrow$ black jack agent

reinforcement: learn with ‘**partial/implicit information**’ (often sequentially)

Mini Summary

Learning with Different Data Label y_n

- **supervised**: all y_n
- unsupervised: no y_n
- semi-supervised: some y_n
- reinforcement: implicit y_n by goodness($\tilde{y}_n$)
  含蓄
- . . . and more!!

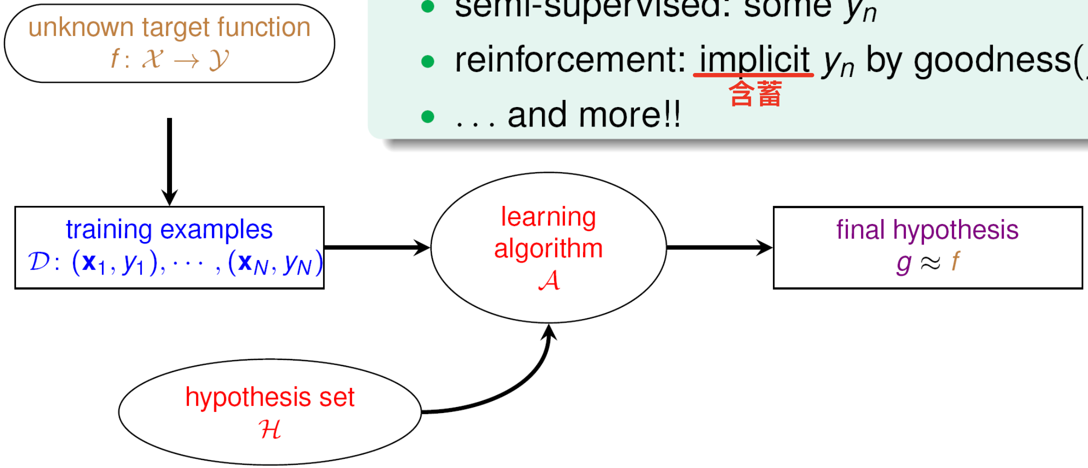

core tool: supervised learning

Fun Time

What is this learning problem?

To build a tree recognition system, a company decides to gather one million of pictures on the Internet. Then, it asks each of the 10 company members to view 100 pictures and record whether each picture contains a tree. The pictures and records are then fed to a learning algorithm to build the system. What type of learning problem does the algorithm need to solve?

1. supervised
2. unsupervised
3. semi-supervised
4. reinforcement

Reference Answer: 3

The $1,000$ records are the labeled $(\mathbf{x}_n, y_n)$; the other $999,000$ pictures are the unlabeled $\mathbf{x}_n$.

Learning with Different Protocol $f \rightarrow (\boldsymbol{x}_n; \boldsymbol{y}_n)$

- Batch Learning: Coin Recognition Revisited

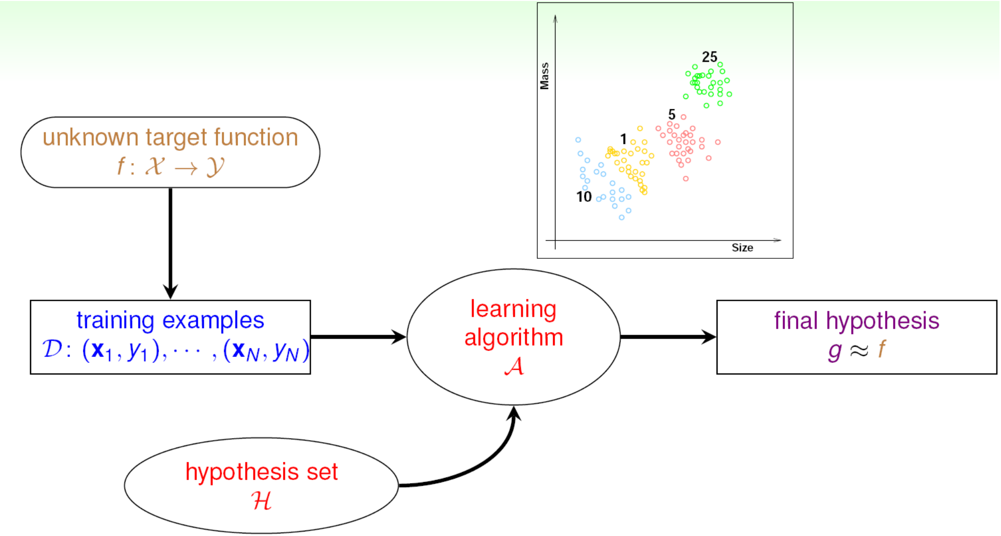

batch supervised multiclass classification:
learn from **all known** data

More Batch Learning Problems

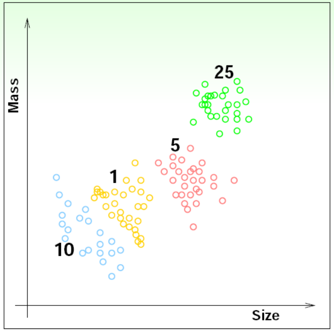

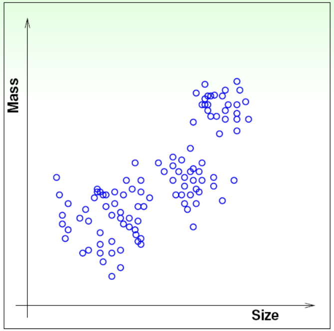

- batch of (email, spam?) ⇒ spam filter
- batch of (patient, cancer) ⇒ cancer classifier
- batch of patient data ⇒ group of patients

batch learning: **a very common protocol**

Online: Spam Filter that ‘Improves’

- batch spam filter:
  learn with known (email, spam?) pairs, and predict with fixed g
- **online** spam filter, which **sequentially**:
  1. observe an email $\mathbf{x}_t$
  2. predict spam status with current $g_t(\mathbf{x}_t)$
  3. receive ‘desired label’ y_t from user, and then update g_t with $(\mathbf{x}_t, y_t)$

Connection to What We Have Learned

- reinforcement learning is often done online (why?)

online: hypothesis ‘improves’ through receiving data instances **sequentially**

Active Learning: Learning by ‘Asking’

Protocol ⇔ Learning Philosophy

- batch: ‘duck feeding’
- online: ‘passive sequential’
- **active: ‘question asking’** (sequentially)
  —query the y_n of the **chosen** $\mathbf{x}_n$

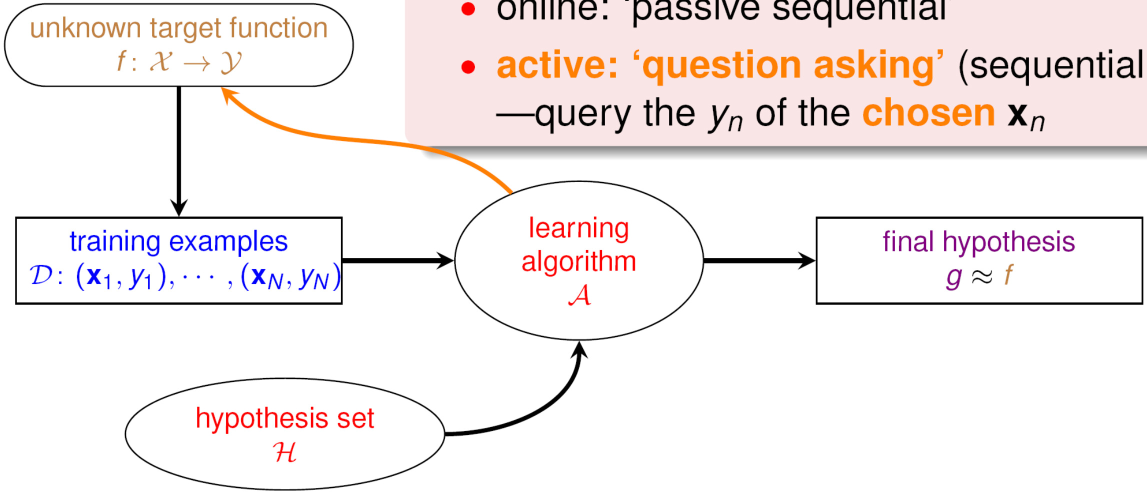

active: improve hypothesis with fewer labels (hopefully) by asking questions **strategically**

Mini Summary

Learning with Different Protocol $f \Rightarrow (\mathbf{x}_n, y_n)$

- **batch**: all known data
- online: sequential (passive) data
- active: strategically-observed data
- ... and more!!

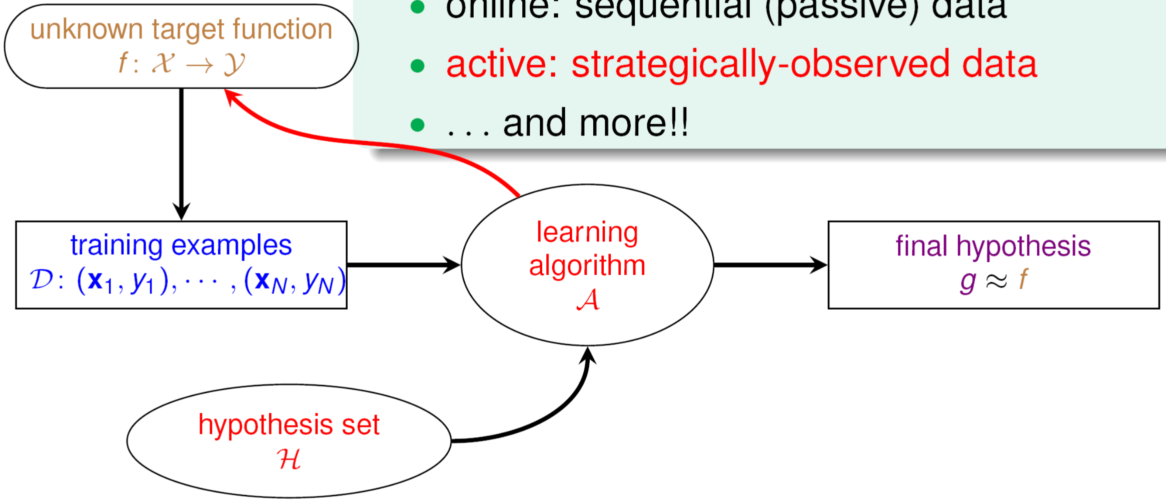

core protocol: batch

Fun Time

What is this learning problem?

A photographer has $100,000$ pictures, each containing one baseball player. He wants to automatically categorize the pictures by its player inside. He starts by categorizing $1,000$ pictures by himself, and then writes an algorithm that tries to categorize the other pictures if it is 'confident' on the category while pausing for (& learning from) human input if not. What protocol best describes the nature of the algorithm?

1. batch
2. online
3. active
4. random

Reference Answer: 3

The algorithm takes a active but naïve strategy: ask when 'confused'. **You should probably do the same when taking a class. :-)**

Learning with Different Input Space X

- Credit Approval Problem Revisited

age	23 years
gender	female
annual salary	NTD 1,000,000
year in residence	1 year
year in job	0.5 year
current debt	200,000

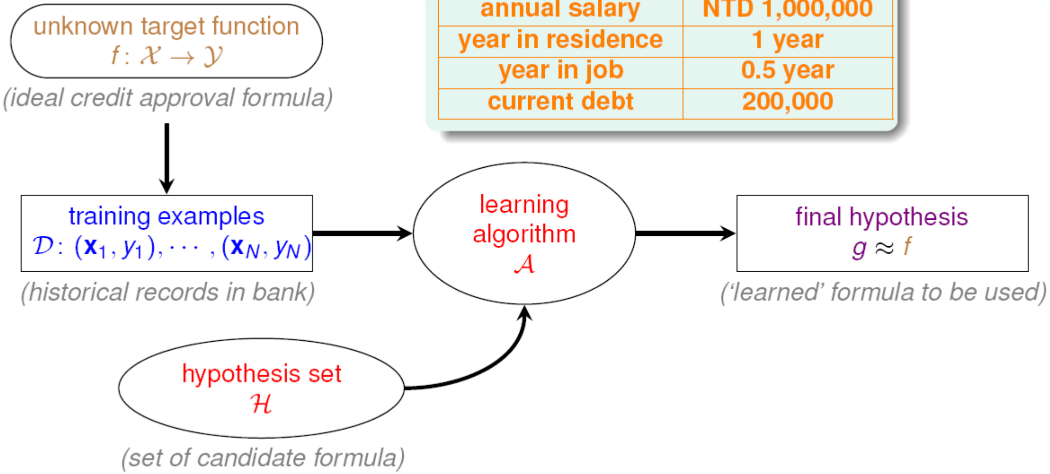

concrete features: each dimension of $\mathcal{X} \subseteq \mathbb{R}^d$ represents 'sophisticated physical meaning'

More on Concrete Features

- 大小 重量 (size, mass) for coin classification
- customer info for credit approval
- patient info for cancer diagnosis
- often including 'human intelligence' on the learning task

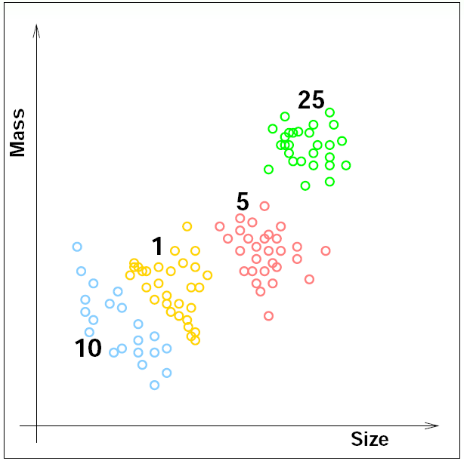

concrete features: the 'easy' ones for ML

Raw Features: Digit Recognition Problem (1/2)

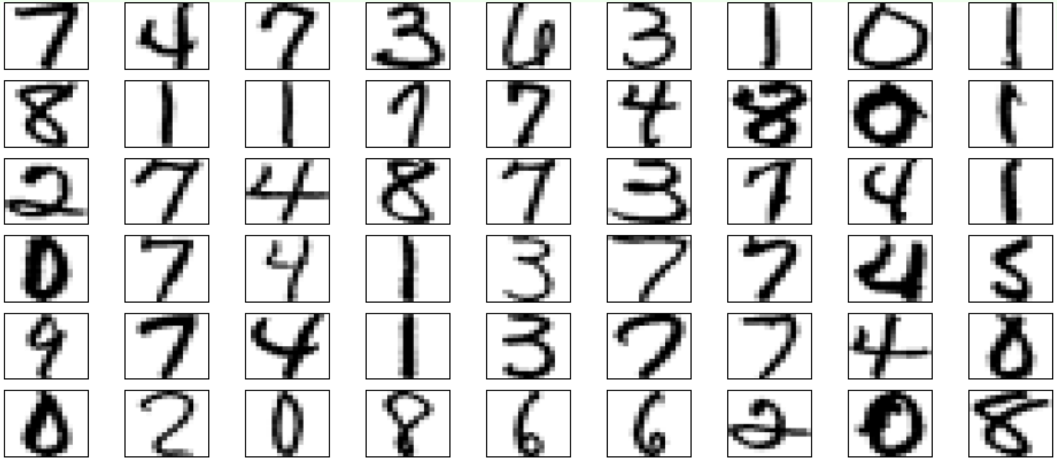

直接使用影像值

- digit recognition problem: features $\Rightarrow$ meaning of digit
- a typical supervised multiclass classification problem

Raw Features: Digit Recognition Problem (2/2)

by Concrete Features

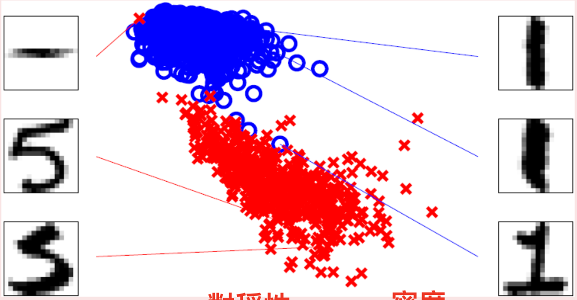

對稱性 密度

$\mathbf{x}$ =(symmetry, density)

by Raw Features

- 16 by 16 gray image $\mathbf{x} \equiv (0, 0, 0.9, 0.6, \cdots) \in \mathbb{R}^{256}$
- '**simple** physical meaning'; thus more difficult for ML than concrete features

Other Problems with Raw Features

- image pixels, speech signal, etc.

raw features: often need human or machines to **convert to concrete ones**

抽象特徵

Abstract Features: Rating Prediction Problem

Rating Prediction Problem (KDDCup 2011)

- given previous (userid, itemid, rating) tuples, predict the rating that some userid would give to itemid?
- a regression problem with $\mathcal{Y} \subseteq \mathbb{R}$ as rating and $\mathcal{X} \subseteq \mathbb{N} \times \mathbb{N}$ **as (userid, itemid)** 喜歡的程度有多高 沒有特別特徵值意
- '**no** physical meaning'; thus even more difficult for ML

Other Problems with Abstract Features

- student ID in online tutoring system (KDDCup 2010)
- advertisement ID in online ad system

abstract: again need '**feature conversion**/extraction/construction'

Mini Summary

Learning with Different Input Space $\mathcal{X}$

- **concrete**: sophisticated (and related) physical meaning
- raw: simple physical meaning
- abstract: no (or little) physical meaning
- ... and more!!

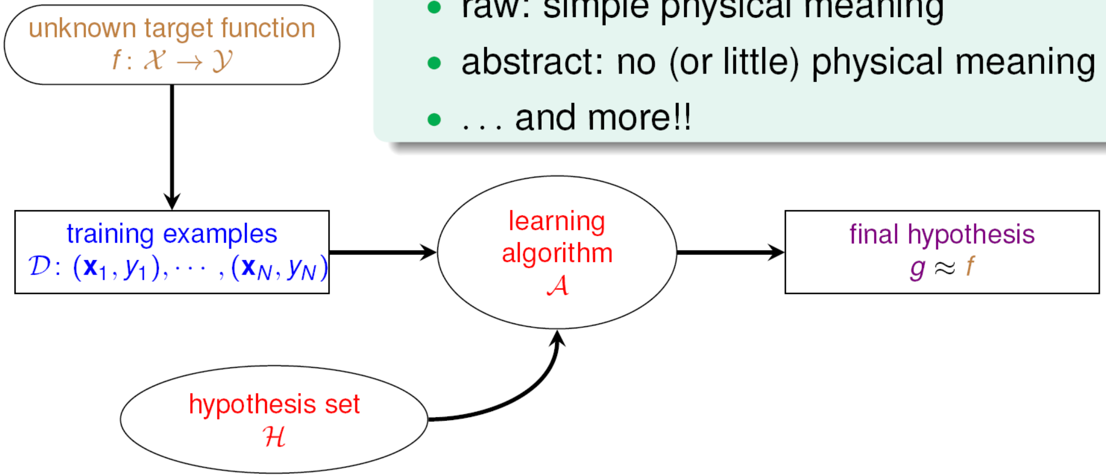

'easy' input: concrete

Summary

- **Types of Learning**

  - Learning with Different Output Space $\mathcal{Y}$

    [classification], [regression], structured
  - Learning with Different Data Label y_n

    [supervised], un/semi-supervised, reinforcement
  - Learning with Different Protocol $f \Rightarrow (\mathbf{x}_n, y_n)$

    [batch], online, active
  - Learning with Different Input Space $\mathcal{X}$

    [concrete], raw, abstract